

# Vadim Uvarov

ML Engineer (production & MLOps) | Data Scientist

PhD in Machine Learning

Male, 33

**Industrial experience:** 8 years

**Location:** Antwerpen, Belgium

**Open to work:** on-site/hybrid in Belgium; remote abroad



ML Engineer with 8 years of experience in production ML, MLOps, and data science across energy, telecom, and AI startup environments. Delivered cloud, on-prem and on-device ML solutions and ETL pipelines, including deployment, monitoring and stakeholder communication.

Having experience both in academy and industry, I can formalize a loosely stated business task, gather inputs from stakeholders and conduct a research & development of a solution either independently or as a part of a team.

I also have an experience of leading a data science team specialized in telecom geodata analysis.

## Skills

### Machine Learning & Deep Learning

- **Classic ML** (classification, regression, clustering, anomaly detection)
- **Time-series forecasting** (energy demand forecasting)
- **Deep learning** (TensorFlow, PyTorch)

### AI Engineering (Gen AI)

- **LLMs & RAG** (Amazon Bedrock)
- Familiar with **LLM**-based applications and actively expanding expertise in this area

### Data analysis & visualization

- Statistics / Hypothesis testing
- Dashboards (**Streamlit**, **simple frontend**)
- Plots: **matplotlib**, **seaborn**, **plotly**

### Python (advanced)

- **Pandas/Polars**, **Scikit-Learn** for data science
- **numpy** / **scipy** for algo implementation
- **pytest** for unit testing
- **commitizen**, **ruff**, **mypy** for clean code

## Contacts

E-mail: [uvarov.vadim42@gmail.com](mailto:uvarov.vadim42@gmail.com)

Phone: +32-489-37-43-99

Links: [LinkedIn](#); [My Website](#)

## Languages

English / C1 (advanced)

Dutch / A2 (elementary and improving)

Russian / C2 (native)

## Visa Status

I'm eligible to work as a freelancer in Belgium (or abroad) and don't require a visa sponsorship.

## Cloud Deployment

- **AWS** (primary), **Kubernetes** (K8s), **Docker**
- **CI/CD**: Gitlab CI/CD, GitHub Actions
- Infrastructure-as-a-code: **Terraform**, **Pulumi**

## MLOps & Production ML

- End-to-end ML lifecycle
- Cloud ML development and deployment (**Amazon Sagemaker**)
- Experiment tracking (**MLflow**)
- Observability: **Grafana**, **Prometheus**
- Model performance tracking, drift detection and retraining strategies

## Big Data

- **Postgres**, **Teradata**, **Hadoop**, **Spark**
- **SQL**: joins, aggregates, window functions

## Software Engineering

- **Git**, unit and integration **testing**, code review, style guide, production-ready code,
- AI-assisted (agentic) coding (**Claude Code CLI**, **GitHub Copilot**)

## Work Experience



### Waiting for a freelance work permit & upskilling

Nov 2025 – Apr 2026

While waiting for a work permit, I've deepened my skills in AWS, ML engineering and AI-assisted coding.



### ML engineer – Centrica [energy company] – Antwerp, Belgium

Apr 2023 – Oct 2025

Centrica Business Solutions helps organizations to develop smart energy solutions.

I worked as a quantitative developer (a blend of a software engineer, data scientist and ML engineer).

- Designed and deployed an **electric vehicle (EV) charging controller**, enabling thousands of UK residents to save on their electricity bills, making it the cheapest way of EV charging in the UK.
- Built production-grade residential **energy demand forecaster** increasing savings by at least 15%.
- Created and deployed a physical-based **solar energy forecasting** model that significantly improved accuracy versus the baseline while preserving interpretability.

My tech stack here was: python, AWS, sklearn, pytorch, Terraform/Pulumi, K8s, Sagemaker, Grafana.



### Data Scientist – Sentiance [AI startup] – Antwerp, Belgium

May 2021 – Jan 2023

Sentiance develops a Mobile SDK which provides insights into users' motion data. I worked on the lifestyle component. It was decided to migrate the solution from backend to the on-device.

- developed a **venue type mapping** model with **Python** and **TensorFlow** for on-device usage to predict venue type visited by user (e.g. shop/restaurant/bar, etc.)
- applied **transfer learning** to improve **venue type mapping** model accuracy
- tracked model training in **MLflow** to support reproducible experiments and ensure quality
- used **TensorFlow (lite)** to develop a multiplatform **on-device** implementation of the algorithm

This allowed for a blazingly fast on-device inference which also increased users' privacy.



### Data Scientist -> Senior Data Scientist – Tele2 Russia [telecom] – Moscow, Russia

Jun 2018 – Apr 2021

At Tele2 telecom company I worked (first as individual contributor and then as a lead) of geodata team.

I've built and deployed several **big data** advanced analytics **ETL pipelines** with using **Spark** and **Airflow**:

- **track processing** for subscriber geolocation using triangulation technique (enhanced accuracy)
- migrated and enhanced **point-of-interest** detection (home/work) pipeline from **Teradata** to **Spark**
- **map matching** solution in **Scala Spark** to snap millions of subscribers' geolocation to a particular street/road segment (here I applied findings from my PhD thesis).

All that enhanced accuracy, reduced time-to market and provided ML features for the other teams.



### Data Scientist – Magnit [food retailer] – Krasnodar, Russia

May 2017 – May 2018

I worked on a **retail site selection** problem: predict the shop potential revenue based on its coordinates.

Another project was **on-shelf availability** prediction system (based on time-series forecasting, not computer vision) to detect and promptly signalize that a particular product should be restocked.

## Education



### Novosibirsk State Technical University – PhD in Machine Learning

2015-2019

My research was related to hidden Markov models. I have defended my PhD thesis in February 2020.



### Novosibirsk State Technical University – MSc in Applied Mathematics and Computer Science

2010-2015

## My projects, articles, speeches

### Pet project: event-driven task queue on AWS

Built an [event-driven task queue on AWS \(Lambda, SQS, Terraform\)](#) to deepen my cloud expertise.  
Repo: <https://github.com/vadim-uvarov/aws-demo-simple-task-queue>.

### Research papers

Recent paper: [“On-Device Deep Learning Location Category Inference Model” by Gadzhi Musaev, Kevin Mets, Rokas Tamosiunas, Vadim Uvarov, Tom De Schepper, Peter Hellinckx \(Belgium, 2022\)](#)

My Research Gate profile: [https://www.researchgate.net/profile/Vadim\\_Uvarov](https://www.researchgate.net/profile/Vadim_Uvarov)

### Medium article “Feature selection using hypothesis testing”

[https://medium.com/@vadim\\_uvarov/feature-selection-using-statistical-testing-2d8e7b5e27b8](https://medium.com/@vadim_uvarov/feature-selection-using-statistical-testing-2d8e7b5e27b8)

### Work Project “Map Matching of Cellular Data using hidden Markov models”

My [conference speech](#) about the technical aspects of this project (in Russian with English subtitles).

[Scientific paper](#) about this project, written in collaboration with my colleagues (in Russian, page 32).

### PhD thesis “Incomplete sequence analysis using hidden Markov models”

Thesis (in Russian): [www.nstu.ru/files/dissertations/dissertaciya\\_1574424031\\_157443103786.pdf](http://www.nstu.ru/files/dissertations/dissertaciya_1574424031_157443103786.pdf)

Video of my thesis defense (in Russian): [link](#). Code and experiments: [link](#).

### Pet project: street signs character recognition (computer vision, deep learning)

I decided to implement a convolutional neural network for classifying images of letters (capital and lowercase) and digits in the wild (Chars74k dataset). I used Python, Theano and Lasagne library for CNN construction, Keras for data augmentation and OpenCV for image preprocessing ([code on GitHub](#)). I achieved 83% accuracy score and 14-th place in [Kaggle training competition](#).

## Recent Courses taken and Books read

Udemy course: [AWS Certified Machine Learning Engineer Associate MLA-c01](#)

Udemy course: [AWS Certified Developer Associate 2025 DVA-C02](#).

Book: “Designing Data-Intensive Applications” by Martin Klepmann.

Book: “Designing Machine Learning Systems” by Chip Huyen.

Coursera specialization (watched lectures): [Deep Learning](#) by Deeplearning.AI.

Coursera specialization: [machine learning and data science](#) by Yandex.

Course (watched lectures): [CS231n](#) CNNs for Visual Recognition by Andrej Karpathy.

I’ve taken Dutch courses, which allowed me to achieve a certified A2 level in Dutch. I’m reading books and practice conversations in Dutch and to further advance my language skills.

## Hobbies

To let my mind unwind after work, I play padel, go to gym and explore Belgium. I play electric guitar and enjoy video games. I like travelling and getting familiar with foreign cultures. I’m an AI enthusiast and use Perplexity for my daily life as well as agentic coding tools (Claude Code CLI) for side projects.